



DEPARTAMENTO DE SEÑALES, SISTEMAS Y RADIOCOMUNICACIONES



# Acoustic Signal Processing

Máster Universitario en Ingeniería de Telecomunicación

Prof. Eduardo López Gonzalo

[eduardo.lopez@upm.es](mailto:eduardo.lopez@upm.es)

Prof. Luis A. Hernández Gómez

[luisalfoso.hernandez@upm.es](mailto:luisalfoso.hernandez@upm.es)

Departamento de Señales, Sistemas y Radiocomunicaciones

E.T.S. Ingenieros de Telecomunicación

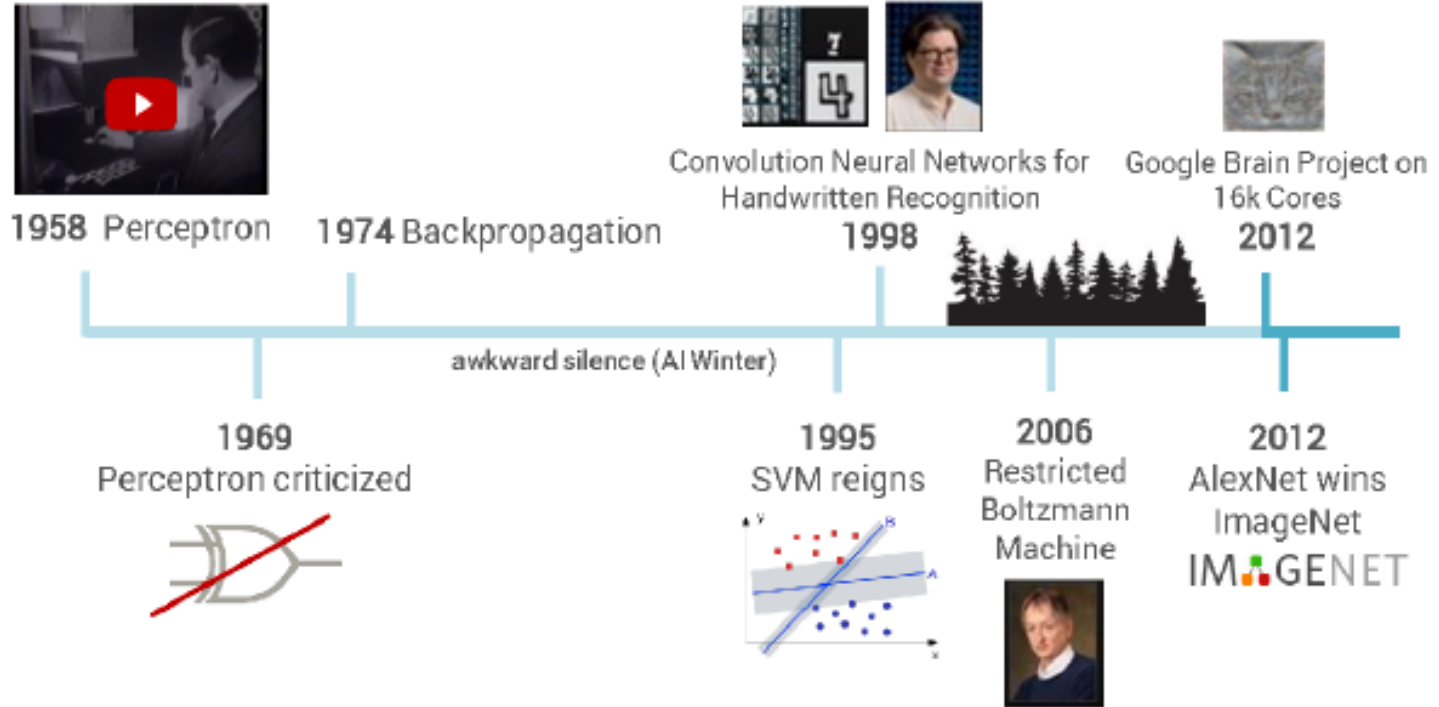
Universidad Politécnica de Madrid

# Deep Learning: Hype or Reality?



# Deep Learning: Hype or Reality?

## A Brief History



Geoffrey Hinton:

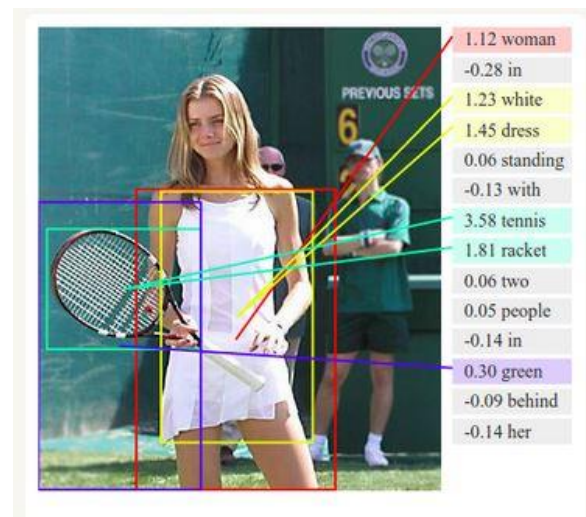
Deep Learning using TensorFlow and TensorFlow-Slim  
Dipendra Jha Northwestern University

# Deep Learning: Hype or Reality?

DNNs better than humans at image recognition



ImageNet  
1000 categories  
1.3M images  
Human error: 5%  
DNN: 3%




Super Human!



## Speech Recognition

According to Microsoft's speech group:





# Deep Learning: Hype or Reality?

## Deep Reinforcement Learning



...going unsupervised! Deep Clustering

# Introduction to Deep Learning but what is new?

## What Changed? Old wine in new bottles



Big Data  
(Digitalization)



Computation  
(Moore's Law, GPUs)




Algorithmic  
Progress

UDACITY

FREE COURSE

Deep Learning  
by Google




# Introduction to Deep Learning

## ...the Machine Learning background...

Learn the whole  
Machine Learning  
context

On line:  
[www.deeplearningbook.org](http://www.deeplearningbook.org)

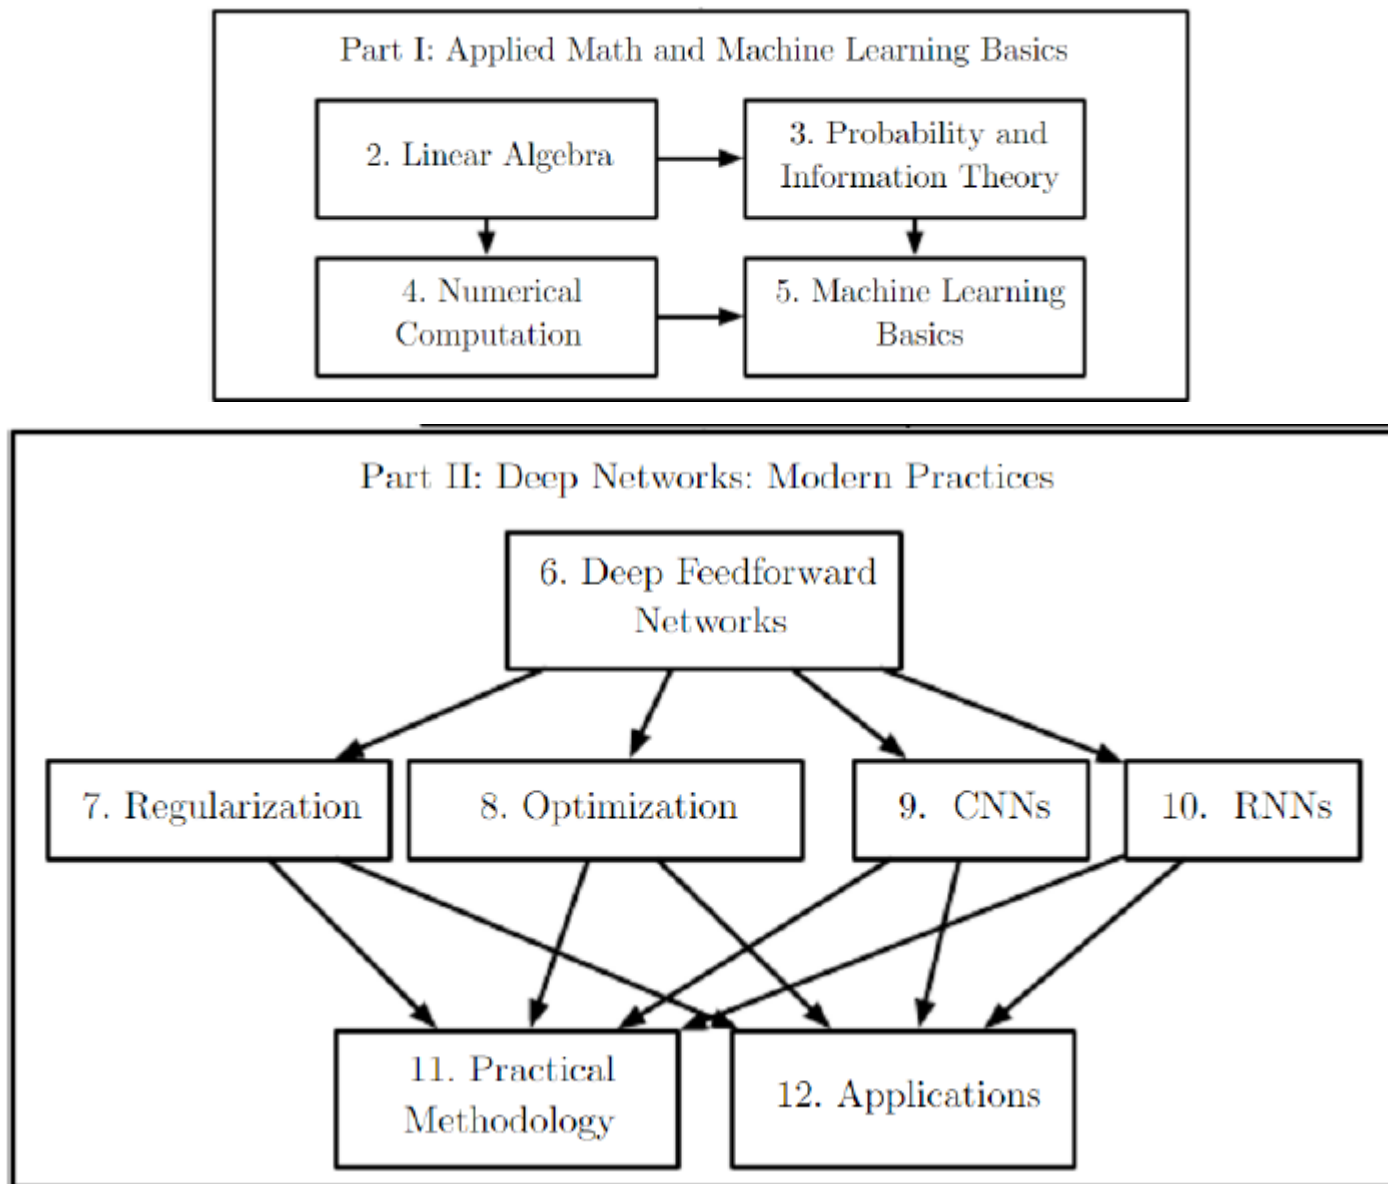


Deep Learning courses

**Prof. Hung-yi Lee** National Taiwan University (NTU) Taipei

# Introduction to Deep Learning


## ...the Machine Learning background...





# Introduction to Deep Learning


## ...the Machine Learning background...




Drivies use case:

*[www.driviesapp.com](http://www.driviesapp.com)*







**Accelerometer**



**Gyroscope**






# From manual - linear classifiers TO Neural Networks

Gyroscopes Energy

$x_2$

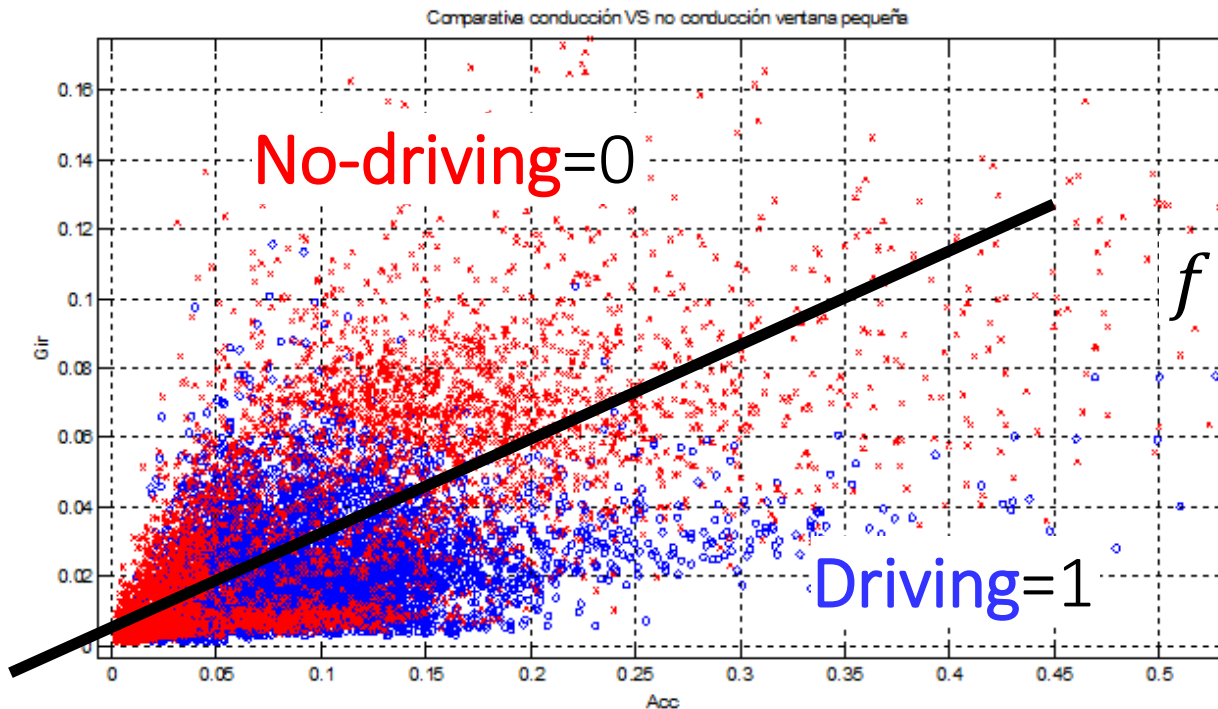


$x_1$  : Accelerometers Energy

Driving detection (**yes/no**)  
= define a **decision function**

Gyroscopes Energy

$x_2$




$$f: R^K \rightarrow R$$

$x_1$  (Accelerometers Energy)

# From linear classifiers TO Neural Networks

## A Linear decision function



## A Linear **decision** function

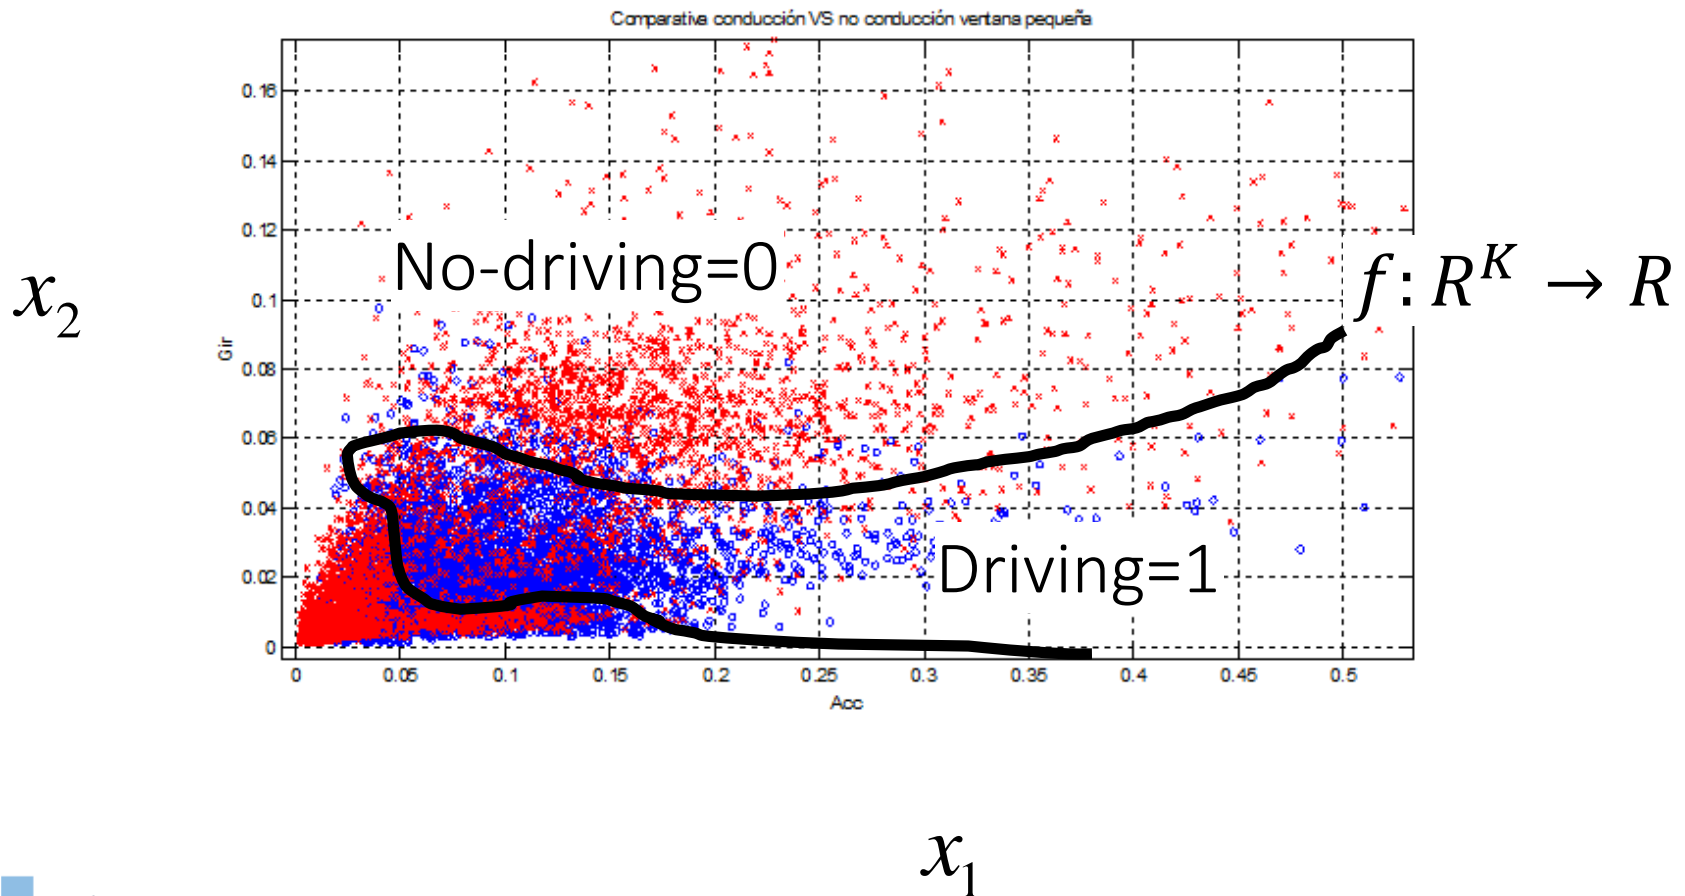
$$y = x_1 w_1 + x_2 w_2 + \cdots + x_K w_K + b$$

$$y = \mathbf{x}^T \mathbf{w}$$



# From linear classifiers TO Neural Networks

## Nonlinear decision function?



# Non-linear decision functions

$$y = x_1 w_1 + x_1^2 w_2 + x_1^3 w_3 + x_1 x_2 w_4 + \cdots + b$$

A linear model of transformed inputs:

$$y = \phi(\mathbf{x})^T \mathbf{w}$$

$\phi(\mathbf{x})$  where  $\phi$  is a non linear transformation


# How choosing the mapping $\phi(.)$ ?

1. To feature engineer  $\phi(.)$
2. Use a very generic  $\phi(.)$  as kernel machines (e.g. SVM, RBF kernel)
3. The strategy of **deep learning** : to learn  $\phi(.)$


# SVM

Hand-crafted  
kernel function

$\phi$



**Input Space**



**Feature Space**

Apply simple  
classifier

From: DL Tutorial  
*Hung-yi Lee*

Source of image: [http://www.gipsa-lab.grenoble-inp.fr/transfert/seminaire/455\\_Kadri2013Gipsa-lab.pdf](http://www.gipsa-lab.grenoble-inp.fr/transfert/seminaire/455_Kadri2013Gipsa-lab.pdf)



# The DL approach: learn $\phi(\cdot)$


Now we have:

- Parameters  $\boldsymbol{\theta}$  that we use to learn  $\phi(\cdot)$  from a broad class of functions
- Parameters  $\mathbf{w}$  that map  $\phi(\mathbf{x})$  to the desired output


$$y = f(\mathbf{x}; \boldsymbol{\theta}, \mathbf{w}) = \phi(\mathbf{x})^T \mathbf{w}$$

# The DL approach: learn $\phi(\mathbf{x})$


...from a broad class of functions



# Neuron approach....



# Neural Network (from Hung-yi Lee “Deep Learning Tutorial”)

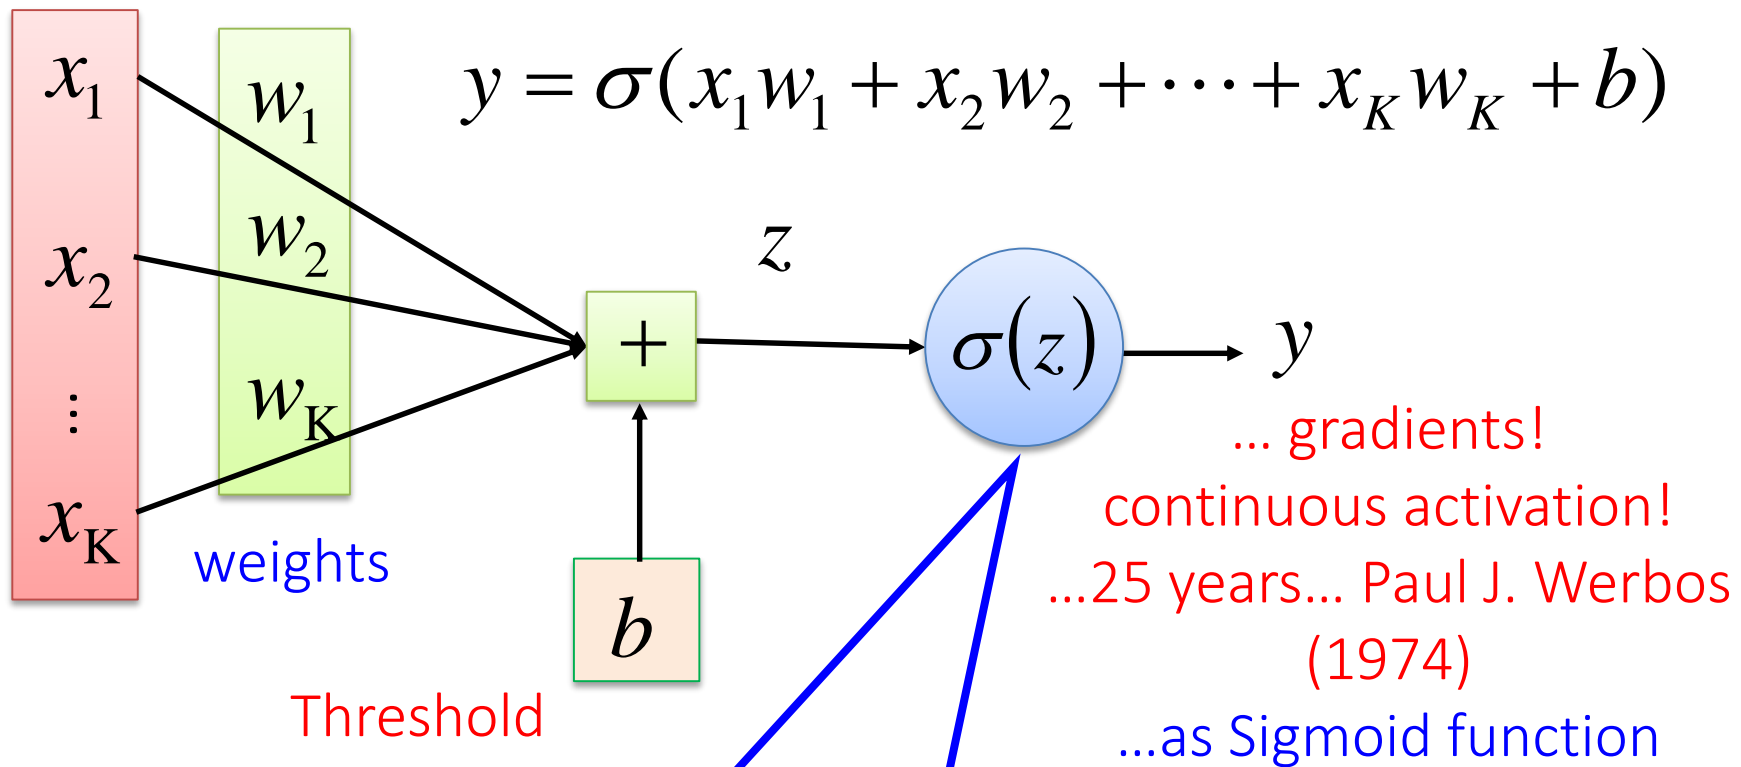


Deep means many hidden layers




# From linear classifiers TO Neural Networks

- Parameters  $\theta$  that we use to learn  $\phi(.)$  from a broad class of functions
- **Weights** and **thresholds** are estimated from training examples:
  - to minimize a **loss function** (i.e. similarity between NN outputs  $y$  and desired outputs  $\hat{y}$  )

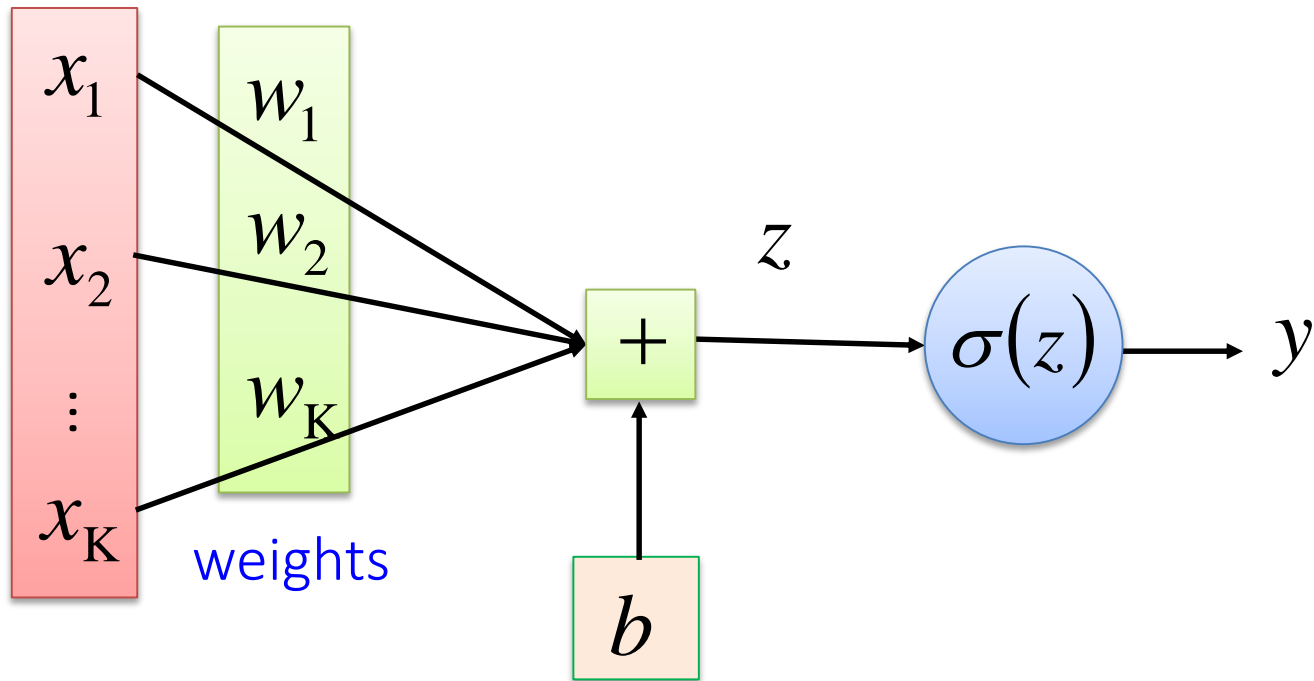


Sigmoid Function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



- But recall that this is also logistic regression!



# From linear classifiers TO Neural Networks

- So let's stop here and start playing with



# Google TensorFlow

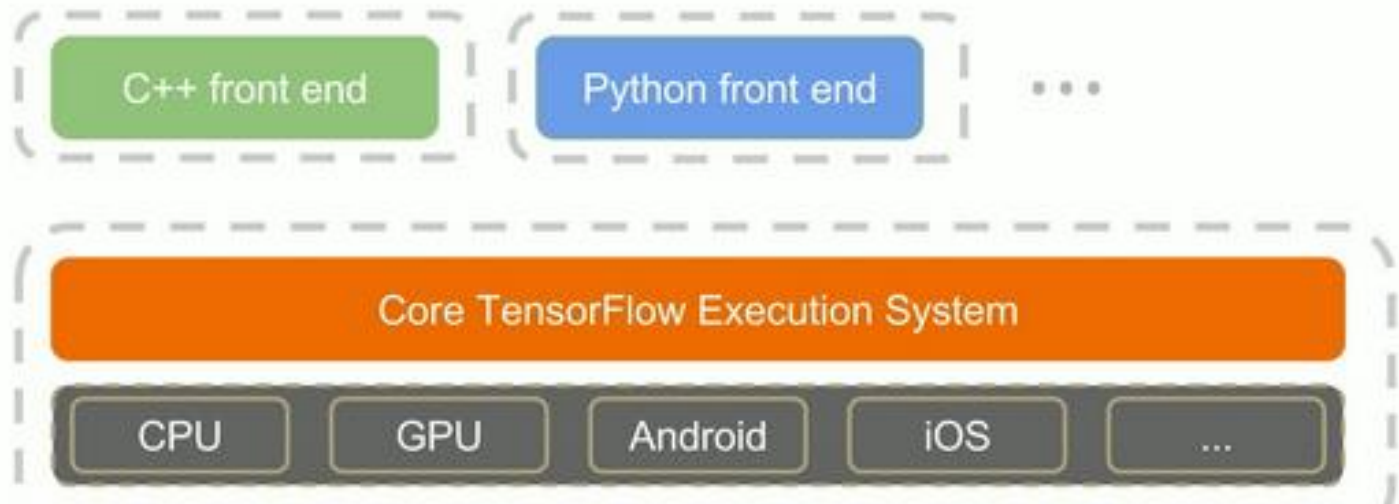


- Library for writing “machine intelligence” algorithms
- Very popular for deep learning and neural networks
- Can also be used for general purpose numerical computations
- Interface in C++ and Python



# TensorFlow: Expressing High-Level ML Computations

- Core in C++
- Different front ends for specifying/driving the computation



A word of caution: the APIs in languages other than Python are not yet covered by the [API stability promises](#).

- [Python](#)
- [C++](#)
- [Java](#)
- [Go](#)

A multidimensional array.



TensorFlow

A graph of operations.


From Mihaela Rosca Talk (Deep Mind)





Generates a computational graph like Theano  
Everything about TensorFlow is here:


<https://www.tensorflow.org>



# TensorFlow Recipe:



- Define a series of expressions
- Initialize variables
- Start a session (launch a graph)
- Run the graph, **feed** some data, **fetch** some values



# TensorFlow Essentials:



**Four types of objects** make TensorFlow unique from other frameworks

- Session
- Computational graph
- Variables
- Placeholder

Let's start playing with



How?

we have prepared some Interactive Python notebooks  
(Jupyter) <http://jupyter.org/about.html>

TSA\_IntroTF\_1.ipynb

[https://github.com/MUIT-TSA/Python/DeepLearning\\_TF\\_Keras/](https://github.com/MUIT-TSA/Python/DeepLearning_TF_Keras/)

TSA\_IntroTF\_1.ipynb

# We recommend you try:

## Colaboratory

It's a Jupyter notebook environment that requires no setup to use and runs entirely in the cloud.

*Colaboratory is free to use.*

# Welcome to Colaboratory!

- Colaboratory is a Google research project created to help disseminate machine learning education and research.
- Colaboratory notebooks are stored in [Google Drive](#) and can

## GPU Support (NEW!)


Colab now supports running TensorFlow computations on a GPU.



# TensorFlow Docs:

“TensorFlow programs are usually structured into a construction phase, that assembles a graph...”

$$y = x_1 w_1 + b$$



```
import tensorflow as tf

x=tf.constant(1.0)
w=tf.constant(6.0)
b=tf.constant(1.5)

y=x*w+b

print(y)
Tensor("add:0", shape=(), dtype=float32)
```

# TensorFlow Docs:

- “A Session object encapsulates the environment in which Tensor objects are evaluated...”
- “..and an execution phase that uses a session to execute ops in the graph”

```
import tensorflow as tf

x=tf.constant(1.0)
w=tf.constant(6.0)
b=tf.constant(1.5)

y=x*w+b

with tf.Session() as sess:
    print(sess.run(y))
```

7.5

HOWEVER *Tensorflow* is open new venues!

## See: Eager Execution for TensorFlow

Announcing TensorFlow 1.5 Friday, January 26, 2018

*With Eager Execution for TensorFlow enabled, you can execute TensorFlow operations immediately as they are called from Python.*


*This makes it easier to get started with TensorFlow, and can make research and development more intuitive.*

# TensorFlow Variables

- “... hold and update **parameters**”

$$y = x_1 w_1 + b$$

- ..and graph (session)  
is executed several times



```
w=tf.Variable(tf.random_uniform([1], -1.0, 1.0))
b=tf.Variable(tf.zeros([1]))
x=tf.constant(1.0)
```

```
# Before starting, initialize the variables
init = tf.global_variables_initializer()
```

```
y=x*w+b
```


```
with tf.Session() as sess:
    sess.run(init)
    for step in range(4):
        print(sess.run(y))
```

# TensorFlow Placeholders



- “... dummy nodes that provide **entry points** to the computational graph”

$$y = x_1 w_1 + b$$



```
W=tf.Variable(tf.random_uniform([1], -1.0, 1.0))
b=tf.Variable(tf.zeros([1]))
x=tf.placeholder(tf.float32)
```

```
# Before starting, initialize the variables
init = tf.global_variables_initializer()
```

```
y=x*w+b
```

```
with tf.Session() as sess:
    sess.run(init)
    for step in range(4):
        print(sess.run(y, feed_dict=(x:1)))
```

# Why TensorFlow?



- Python + Numpy
- Graph based, easy to model
- Faster compile times than Theano
- **Tensorboard** for Visualization
- **<http://playground.tensorflow.org>**
- Open Sourced
- Data and Model Parallelism
- Distributed supported

# But what's a Tensor?



## Tensor Ranks, Shapes, and Types

Briefly: ***A tensor is an array of  $n$ -dimension containing the same type of data***

- **Tensor rank** (sometimes referred to as order or degree or  $n$ -dimension) is the number of dimensions of the tensor.

Rank	Math entity	Python example
0	Scalar (magnitude only)	<code>s = 483</code>
1	Vector (magnitude and direction)	<code>v = [1.1, 2.2, 3.3]</code>
2	Matrix (table of numbers)	<code>m = [[1, 2, 3], [4, 5, 6], [7, 8, 9]]</code>
3	3-Tensor (cube of numbers)	<code>t = [[[2], [4], [6]], [[8], [10], [12]], [[14], [16], [18]]]</code>
$n$	$n$ -Tensor (you get the idea)	<code>....</code>



# But what's a Tensor?



## Tensor Ranks, Shapes, and Types

- **Tensor shape:** The TensorFlow documentation uses three notational conventions to describe tensor dimensionality: rank, shape, and dimension number.

Rank	Shape	Dimension number	Example
0	<code>[]</code>	0-D	A 0-D tensor. A scalar.
1	<code>[D0]</code>	1-D	A 1-D tensor with shape <code>[5]</code> .
2	<code>[D0, D1]</code>	2-D	A 2-D tensor with shape <code>[3, 4]</code> .
3	<code>[D0, D1, D2]</code>	3-D	A 3-D tensor with shape <code>[1, 4, 3]</code> .
n	<code>[D0, D1, ... Dn-1]</code>	n-D	A tensor with shape <code>[D0, D1, ... Dn-1]</code> .

Shapes can be represented via Python lists / tuples of ints, or with the `tf.TensorShape`.

# But what's a Tensor?



## Tensor Ranks, Shapes, and Types

Data type	Python type	Description
DT_FLOAT	<code>tf.float32</code>	32 bits floating point.
DT_DOUBLE	<code>tf.float64</code>	64 bits floating point.
DT_INT8	<code>tf.int8</code>	8 bits signed integer.
DT_INT16	<code>tf.int16</code>	16 bits signed integer.
DT_INT32	<code>tf.int32</code>	32 bits signed integer.
DT_INT64	<code>tf.int64</code>	64 bits signed integer.
DT_UINT8	<code>tf.uint8</code>	8 bits unsigned integer.
DT_UINT16	<code>tf.uint16</code>	16 bits unsigned integer.
DT_STRING	<code>tf.string</code>	Variable length byte arrays. Each element of a Tensor is a byte array.
DT_BOOL	<code>tf.bool</code>	Boolean.

.....

.....

# But what's a Tensor?



A Tensor in TensorFlow has 2 shapes! The **static shape**  
AND the **dynamic shape**

- The static shape can be read using the `tf.Tensor.get_shape()` method: this shape is inferred from the operations that were used to create the tensor, and may be partially complete.

```
x = tf.placeholder(tf.int32, shape=[4])
print x.get_shape()
# ==> '(4,)'
```

- If the static shape is not fully defined, the dynamic shape of a Tensor `t` can be determined by evaluating `tf.shape(t)`.

```
y, _ = tf.unique(x)
print y.get_shape()
# ==> '(?,)'
```

# But what's a Tensor?



A Tensor in TensorFlow has 2 shapes! The **static shape**  
AND the **dynamic shape**

- getting dynamic shape of a Tensor `t` can be determined by evaluating `tf.shape(t)`.

```
y, _ = tf.unique(x)
```

```
sess = tf.Session()  
print sess.run(y, feed_dict={x: [0, 1, 2, 3]}).shape  
# ==> '(4,)'
```

```
print sess.run(y, feed_dict={x: [0, 0, 0, 0]}).shape  
# ==> '(1,)'
```

# But what's a Tensor?



mathematical operations to manipulate the *tensors*

Operation	Description
<code>tf.add</code>	sum
<code>tf.sub</code>	subtraction
<code>tf.mul</code>	multiplication
<code>tf.div</code>	division
<code>tf.mod</code>	module
<code>tf.abs</code>	return the absolute value
<code>tf.neg</code>	return negative value
<code>tf.sign</code>	return the sign
<code>tf.inv</code>	returns the inverse
<code>tf.square</code>	calculates the square
<code>tf.round</code>	returns the nearest integer
<code>tf.sqrt</code>	calculates the square root
<code>tf.pow</code>	calculates the power
<code>tf.exp</code>	calculates the exponential
<code>tf.log</code>	calculates the logarithm

.....

# But what's a Tensor?



*Dealing with shapes is one of the major issues when working with TensorFlow!*

# Placeholders: Feeding data...



`tf.placeholder()`

*Placeholders allow you to pass in numpy arrays of data*


At run time, we feed using a **feed\_dic**

Arguments: data type (mandatory), shape (optional),  
name (optional)

# Let's try a simple linear classifier

$$y = x_1 w_1 + x_2 w_2 + b$$


Inputs







# Binary (two-class) classifier from synthetic data






TensorFlow

# Binary (two-class) classifier from synthetic data

**Training:** How to find the parameters  $\theta : W, b$  ?



**OPTIMIZATION =>**

Classification Error  
Cost a function of the parameters =  $C(\theta)$


TSA\_IntroTF\_2.ipynb



For particular values of  $W$  and  $b$ :  $x_1 w_1 + x_2 w_2 + b = 0$



Red line is the set of points defined by:  
... arbitrary ....

$$x_2 + 1.23 * x_1 - 0.55 = 0$$



Green points will give:  $x_2 + 1.23 * x_1 - 0.55 > 0$

Blue points will give:  $x_2 + 1.23 * x_1 - 0.55 < 0$







# What Cost Function ?

- Classification error (number of errors)?  
How derivate? gradient?
- Better think on terms of **probability** & look for nice ways to derivate...

# From scores to probabilities:

- Binary classification is easy: use sigmoid or logistic function



# Cost function:

How to measure accuracy?

## Cross-entropy and Maximum Likelihood Estimation



# Intuitive approach for a two-class (binary) classifier

Take a data set:

$\{x_n, l_n\}$  of  $N$  vectors  $x_n$  belonging to two classes (labels  $l_n = 0, 1$ )

Consider our classifier as an estimator of the conditional probability:

- $p(l_n = 1|x_n) = p_{1n}$
- ...and so...  $p(l_n = 0|x_n) = 1 - p_{1n}$

# A PERFECT classifier

- If  $l_n=1$  then  $p(l_n = 1|x_n) = p_{1n}=1$
- If  $l_n=0$  then  $p(l_n = 1|x_n) = p_{1n}=0$

In a single expression:

$$p_{1n}^{l_n} (1 - p_{1n})^{(1-l_n)}$$

In a perfect classifier this must be:  
a “probability” always 1 for every data  $n$ !!!

# A PERFECT classifier

All the probabilities/likelihood for each sample should be 1

Maximum Likelihood estimation of  $\theta$

$$\theta_{ML} = \prod_{n=1}^N p_{1n}^{l_n} (1 - p_{1n})^{(1-l_n)}$$

From that a Cost Function can be obtained taking: the negative logarithm = Cross-entropy

$$\text{Cross entropy} = - \sum_{n=1}^N (l_n \log(p_{1n}) + (1 - l_n) \log(1 - p_{1n}))$$

# Cross entropy loss function

From Maximum Likelihood to Minimum Cross-entropy

$$\text{Cross entropy} = - \sum_{n=1}^N (l_n \log(p_{1n}) + (1 - l_n) \log(1 - p_{1n}))$$

Shannon entropy  $H(\cdot)$  : of uncertainty in a probability distribution  $P$

$$H(x) = -E_{x \sim P} [\log(P(x))]$$

Cross-entropy (two distributions :  $P_{data}$  and  $P_{model}$ )

$$-E_{x \sim P_{data}} [\log(P_{model}(x))]$$

# Cross entropy loss function

ANOTHER “PRACTICAL” POINT OF VIEW:

$$\text{Cross entropy} = - \sum_{n=1}^N (l_n \log(p_{1n}) + (1 - l_n) \log(1 - p_{1n}))$$

**HEAVILY PENALIZES** gross errors:



- $l_n = 1$  and  $p_{1n}$  close to 0  **$\log(0)!!$**
- $l_n = 0$  and  $p_{1n}$  close to 1 !! as  $(1 - p_{1n})$  close to 0  **$\log(0)!!$**

Now let's feed our Training data to a simple linear classifier

$$x_2 + 1.23 * x_1 - 0.55 = z$$

Inputs

Train data



- Probabilities
- Cross entropy



```
W=tf.constant([[1.23], [1.0]],name="weights")  
b=tf.constant(-0.55,name="bias")
```



## Inputs

$$x_2 + 1.23 * x_1 - 0.55 = z$$

```
X = tf.placeholder("float", shape=[None, 2])
```

## Train data

```
labels = tf.placeholder("float", shape=[None])
```



# Predictor is now the logistic function

```
pred = tf.sigmoid(tf.to_double(tf.reduce_sum(tf.matmul(X, W),  
axis=[1]) + b))
```

# Cost function is cross-entropy

```
cost = -tf.reduce_sum(tf.to_double(labels) * tf.log(pred) + (1-  
tf.to_double(labels)) * tf.log(1-pred))
```



Probabilities  
Cross entropy



TensorFlow

## Inputs

## Train data



with tf.Session() as sess:

```
sess.run(init)
```

```
pred, cost=sess.run(pred, cost, feed_dict={X:  
train_X})
```

- Probabilities
- Cross entropy





TensorFlow

Check that our results are  
the same as before

# How to get the best $W$ and $b$ values?

(that is: those who give you the lowest cost)

`GradientDescentOptimizer(learning_rate).minimize(cost)`

Next Slides are from:  
Deep Learning Tutorial

李宏毅


Hung-yi Lee

# OPTIMIZATION

## Gradient Descent

Cost  
Or  
Loss function

$$C(\theta)$$



Assume there are only two parameters  $w_1$  and  $w_2$  in a network.

$$\theta = \{w_1, w_2\}$$

# OPTIMIZATION

## Gradient Descent

Assume there are only two parameters  $w_1$  and  $w_2$  in a network.

$$\theta = \{w_1, w_2\}$$

Randomly pick a starting point  $\theta^0$


Compute the negative gradient at  $\theta^0$

$$\rightarrow -\nabla C(\theta^0)$$


Times the learning rate  $\eta$

$$\rightarrow -\eta \nabla C(\theta^0)$$


Error Surface




# Choosing Step Size



Too small: converge very slowly



Too big: overshoot and even diverge



Reduce size over time

Theoretical convergence results for various step sizes

A common step size is  $\alpha_i = \frac{\alpha}{n\sqrt{i}}$


# Training Points
 $\alpha$  — Constant
Iteration #

Source:



BerkeleyX: CS190.1x Scalable Machine Learning

# Gradient Descent



Randomly pick a starting point  $\theta^0$

Compute the negative gradient at  $\theta^0$

➡  $-\nabla C(\theta^0)$


Times the learning rate  $\eta$

➡  $-\eta \nabla C(\theta^0)$



# Local Minima

- Gradient descent never guarantee global minima



Different initial  
point  $\theta^0$




Reach different minima,  
so different results

Who is Afraid of Non-Convex  
Loss Functions?

[http://videolectures.net/eml07\\_lecun\\_wia/](http://videolectures.net/eml07_lecun_wia/)




## Besides local minima .....




# In physical world .....

- Momentum



# Momentum

Still not guarantee reaching global minima, but give some hope .....



# So now let's train!:



- **W** and **b** will be variables (tensors)

```
W = tf.Variable(tf.zeros([2, 1], "float"), name="weight")
```

```
b = tf.Variable(tf.zeros([1], "float"), name="bias")
```

- **and we use cross-entropy and gradient descend**

```
cost = -tf.reduce_sum(tf.to_double(labels) * tf.log(pred) + (1-tf.to_double(labels))  
* tf.log(1-pred))
```

```
# Gradient descent
```

```
learning_rate = 0.001
```

```
optimizer = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
```

- We have used all the training data several runs (Epochs)



```
with tf.Session() as sess:  
    sess.run(init)  
  
    # We can Run the optimization algorithm several times  
    for i in range(100):  
        cost_out,W_out,b_out,_=sess.run([cost, W,b, optimizer], feed_dict={X:  
train_X, labels: train_labels})  
        print("Epoch : %d Cost= %s"%(i,cost_out))  
        print(W_out)  
        print(b_out)
```

# When large amounts of data divide data into **mini-batches**



- Most optimization algorithms converge much faster (in terms of total computation, not in terms of number of updates) if they are allowed to rapidly compute approximate estimates of the gradient rather than slowly computing the exact gradient.
- Another consideration motivating statistical estimation of the gradient from a small number of samples is redundancy in the training set.

- Optimization algorithms that use the entire training set are called **batch or deterministic gradient**
- **Stochastic Gradient Descend (SGD):** Optimization algorithms that use only a single example at a time
- Most algorithms used for deep learning fall somewhere in between: **minibatch or minibatch stochastic methods**

## Confusing terminology:

- The word “batch” is also often used to describe the minibatch used by minibatch stochastic gradient descent.
- It is very common to use the term “batch size” to describe the size of a minibatch

See more details on how choosing minibatch size on Deep Learning Book (Chap 8 : Optimization)



.... See details in Notebook  
... and practice with it....

- Random initialization of variables
- Stepsize
- Optimizers
- Interactive Session
- Tf Debugging?

# Deep Learning Seminar (materials)

Deep Learning using TensorFlow and TensorFlow-Slim

Dipendra Jha Northwestern University

[dipendra009@gmail.com](mailto:dipendra009@gmail.com) <https://www.linkedin.com/in/dipendra009>

Deep Learning courses

Prof. Hung-yi Lee National Taiwan University (NTU) Taipei

Introduction to Deep Learning

[Yingyu Liang](#) Princeton University

<http://jrmeyer.github.io/tutorial/2016/02/01/TensorFlow-Tutorial.html>

<http://www.psi.toronto.edu/~jimmy/ece521/Tut1.pdf>

...and of course look videos/course by **Geoffrey Hinton**: The Godfather of Deep Learning